

CM124 Spring 2019 - Programming Assignment

(due March 13th)

January 31, 2019

Haplotype phasing in recently admixed populations

In this assignment you are going to develop and implement an algorithm that takes as an input the genotypes of individuals coming from admixed populations and infers the haplotypes of the individuals.

Genotyping technologies provide us with strings over $\{0, 1, 2\}$, representing the number of copies of the reference allele at each SNP. However, these technologies do not allow us to measure complete haplotypes (i.e. what is exactly the sequence of alleles in each of the two chromosomes of an individual). The problem gets further complicated when phasing haplotypes of individuals coming from recently admixed populations.

Recall that the creation of an **admixed population** is a process that initially starts with two or more **homogeneous populations** (i.e. each individual is strictly coming from one particular population). Then, at each new generation, following a process of **random mating**, each individual gets two haplotypes, one per parent, from the previous populations (the two haplotypes may come from parents who are part of different populations). Importantly, each of the two haplotypes is going through some recombinations with an unknown recombination rate r . In our case, we consider admixed individuals from two recently admixed populations (i.e. not many generations have passed since the homogeneous populations), under the assumption of random mating.

Develop and implement an algorithm for phasing haplotypes of individuals coming from recently admixed populations. You are free to use any of the methodologies that we have seen in class or that you can find online, and you can use existing software packages as long as they were not designed for haplotype phasing.

Assignment files and compiling your solutions

This assignment is accompanied by the following files:

1. EXAMPLE GENOTYPES: “example_data_1.txt”, “example_data_2.txt”, “example_data_3.txt” - three example datasets to develop and test your method on. Each file contains the **genotypes** for a set of

individuals. Format of files: each row in the file is a genomic position (SNP), each column in the file is an individual. Values are separated by spaces.

2. EXAMPLE HAPLOTYPES: “example_data_1_sol.txt”, “example_data_2_sol.txt”, “example_data_3_sol.txt”
- each file contains the **true haplotypes** for ones of the three sets of example genotypes. Format of files: each row in the file is a genomic position (SNP), each pair of consecutive columns in the file is an individual (i.e. for an individual’s genotype column i , the corresponding haplotypes are located at column $2i$ and $2i + 1$). Values are separated by spaces.
3. TEST GENOTYPES: “test_data_1.txt”, “test_data_2.txt” - each file contains the genotypes for a set of individuals. Once you have developed your method, run your method on these test genotypes and output your estimated haplotypes. These are the haplotypes you will submit for grading. Format of files: each row in the file is a genomic position (SNP), each column in the file is an individual. Values are separated by spaces.
4. GENOTYPES POSITIONS: “example_data_1_geno_positions.txt“, “example_data_2_geno_positions.txt“, “example_data_3_geno_positions.txt“, “test_data_1_geno_positions.txt“, and “test_data_2_geno_positions.txt“
- for each of the example and test dataset, you get the physical positions of the genotypes on the chromosome. This additional information can be incorporated into your algorithm (however, you are not required to use it).

In your final submission you are required to submit your estimated haplotypes for all the individuals in the files “test_data_1.txt” and “test_data_2.txt”. Specifically, submit two files “test_data_1_sol.txt” and “test_data_2_sol.txt”, each with the inferred haplotypes of the genotypes in the corresponding test file using the same format as the example haplotype files (i.e. “example_data_1_sol.txt”, see above).

Importantly, the example datasets are merely provided to you as examples of possible input/output. Note that the input of these datasets do not necessarily represent the input of the test data.

Evaluation of performance

The metric we will be using to measure the performance is the switch accuracy, which is formally defined as $\frac{(n-1-sw)}{(n-1)}$, where n denotes the number of heterozygous sites and sw is the number of switches between neighboring heterozygous sites needed in the computer-phased haplotype to recover the original haplotype sequence (see Lin et al. *AJHG* 2002 for a pictorial explanation).

We are providing an evaluation script you can use to evaluate your algorithm. The same script will be used to grade your final homework submissions. This script requires the *R* software environment, and can be run using the following command:

```
Rscript calculate_switch_accuracy.R [estimated haplotypes file] [true haplotypes file]
```

One of the test files will be randomly selected for use in the fortnightly competitions (see below), and the other test file will be used for the final assignment grade.

General instructions:

- Submissions are accepted for groups of 1 to 3 people
- Your solution will be graded as follows:
 - Report and Code (30 points): Submit your code and a brief summary with a description of your algorithm and implementation (one page maximum). Please submit your report as a PDF document. In addition, prepare a file “readme.txt” which includes the command line commands which are to be executed in order to run your code on the three test datasets.
 - Performance (70 points): algorithm performance will be graded in 3 ways (cumulative):
 - * (20 points): the switch accuracy of your algorithm is better than an algorithm that chooses randomly (i.e. better than 0.5).
 - * (20 points): the switch accuracy of your algorithm is better than 0.7.
 - * (30 points): your algorithm’s performance compared to other groups. This will be calculated as follows: $30 \times \left(1 - \frac{\text{your rank} - 1}{\text{\#groups}}\right)$. For example, if your performance is ranked number 5 in the class and there are 40 groups then your score will be 27. If there are ties in the ranking between two groups, the average rank will be used.
 - * (-5 Points): For every additional hour of run-time. This means that your code should run in less than 1 hour on a modern single-core machine. This is needed to make marking feasible at the end of the quarter.
- For this assignment you are free to use whichever **script programming language** you prefer (i.e. a language which does not require compilation, such as R, Python, etc.). Note that your code will be tested on a Linux/Unix machine.

Fortnightly competitions

Throughout the quarter we will hold fortnightly competitions through CCLE so that you can check your progress on the test datasets and see how others in the class are going.

- Submit your solutions electronically via CCLE by submitting a single zip file. Your zip file should contain two files “test_data_1_sol.txt” and “test_data_2_sol.txt”, each with the inferred haplotypes of the genotypes in the corresponding test file using the same format as the example haplotype files (i.e. “example_data_1_sol.txt”, see above).

- Competition dates are below. Note: CCLE submission will only be available approximately 48 hours before it closes in the folder for this weeks class.
 - Friday, the 8th of February (Week 5).
 - Friday, the 22nd of February(Week 7).
 - Friday, the 8th of March (Week 9).
- The title of the zip file should be your team name (e.g., “hphasers.zip“) and your zip file should only contain the two files discussed above. Note: this differs from the content of your zip files for the final submission.

Final Submission (due March 13th)

- Submit your final solutions electronically via CCLE by submitting a single zip file. Your zip file should include:
 - Your solutions for the two test datasets (two files)
 - Summary/report file
 - Code file(s)
 - “readme.txt” with the execution instructions (see above)

The title of the zip file should be the UID of the student (e.g., “123456789.zip”). If you submit together with other students, separate your UID numbers by a dot (e.g., “123456789.987654321.zip”).